

Three hour tutorial

data.table

30 June 2014

Matt Dowle

Overview

- | | |
|--------------------------------|---------|
| ■ data.table in a nutshell | 20 mins |
| ■ Client server recorded demo | 20 mins |
| ■ Main features in more detail | 2 hours |
| ■ Q&A | 20 mins |

What is data.table?

- Think `data.frame`, inherits from it
- `data.table()` and `?data.table`

Goals:

- Reduce programming time
 - fewer function calls, less variable name repetition
- Reduce compute time
 - fast aggregation, update by reference
- In-memory only, 64bit and 100GB routine
- Useful in finance but wider use in mind, too
 - e.g. genomics

Reducing programming time

```
trades[  
  filledShares < orderedShares,  
  sum( (orderedShares-filledShares)  
        * orderPrice / fx ),  
  by = "date,region,algo"  
]
```

R	:	i	j	by
---	---	---	---	----

SQL	:	WHERE	SELECT	GROUP BY
-----	---	-------	--------	----------

Reducing compute time

e.g. 10 million rows x 3 columns x,y,v 230MB

```
DF[DF$x=="R" & DF$y==123,] # 8 s
```

```
DT[.("R",123)] # 0.008s
```

```
tapply(DF$v,DF$x,sum) # 22 s
```

```
DT[,sum(v),by=x] # 0.83s
```

See above in timings vignette (copy and paste)

Fast and friendly file reading

e.g. 50MB .csv, 1 million rows x 6 columns

`read.csv("test.csv")` # 30-60s

`read.csv("test.csv", colClasses=,
 nrows=, etc...)` # 10s

`fread("test.csv")` # 3s

e.g. 20GB .csv, 200 million rows x 16 columns

`read.csv("big.csv", ...)` # hours

`fread("big.csv")` # 8m

Update by reference using :=

Add new column "sectorMCAP" by group :

```
DT[,sectorMCAP:=sum(MCAP),by=Sector]
```

Delete a column (0.00s even on 20GB table) :

```
DT[,colToDelete:=NULL]
```

Be explicit to really copy entire 20GB :

```
DT2 = copy(DT)
```

Why R?

- 1) R's lazy evaluation enables the syntax :
 - `DT[filledShares < orderedShares]`
 - query optimization before evaluation
- 2) Pass DT to any package taking DF. It works.
`is.data.frame(DT) == TRUE`
- 3) CRAN (cross platform release, quality control)
- 4) Thousands of statistical packages to use with `data.table`

Client/server recorded demo

<http://www.youtube.com/watch?v=rvT8XThGA8o>

Main features in more detail ...

Essential!

- Given a 10,000 x 10,000 matrix in any language
- Sum the rows
- Sum the columns
- Is one way faster, and why?

setkey(DT, colA, colB)

- Sorts the table by colA then colB. That's all.
- Like a telephone number directory: last name then first name
- X[Y] is just binary search to X's key
- You **DO** need a key for joins X[Y]
- You **DO NOT** need a key for by= (but many examples online include it)

Joins: X[Y]

- Vector search vs binary search
- One column == is ok, but not 2+ due to temporary logicals (e.g. slide 5 earlier)
- J(), .(), list(), data.table()
- CJ()
- SJ()
- nomatch
- mult

"Cold" by (i.e. without setkey)

Consecutive calls unrelated to key are fine and common practice :

```
> DT[, sum(v), by="x,y"]  
> DT[, sum(v), by="z"]  
> DT[, sum(v), by=colA%%5]
```

Also known as "ad hoc by"

DT[i, j, by]

- Out loud: "Take **DT**, subset rows using **i**, then calculate **j** grouped by **by**"
- Once you grok the above reading, you don't need to memorize any other functions as all operations follow the same intuition as base.



3



I have a data frame that is some 35,000 rows, by 7 columns. it looks like this:

```
head(nuc)
```

	chr	feature	start	end	gene_id	pctAT	pctGC	length
1	1	CDS	67000042	67000051	NM_032291	0.600000	0.400000	10
2	1	CDS	67091530	67091593	NM_032291	0.609375	0.390625	64
3	1	CDS	67098753	67098777	NM_032291	0.600000	0.400000	25
4	1	CDS	67101627	67101698	NM_032291	0.472222	0.527778	72
5	1	CDS	67105460	67105516	NM_032291	0.631579	0.368421	57
6	1	CDS	67108493	67108547	NM_032291	0.436364	0.563636	55

gene_id is a factor, that has about 3,500 unique levels. I want to, for each level of gene_id get the min(start), max(end), mean(pctAT), mean(pctGC), and sum(length).

I tried using lapply and do.call for this, but it's taking forever +30 minutes to run. the code I'm using is:

```
nuc_prof = lapply(levels(nuc$gene_id), function(gene){  
  t = nuc[nuc$gene_id==gene, ]  
  return(list(gene_id=gene, start=min(t$start), end=max(t$end), pctGC =  
    mean(t$pctGC), pct = mean(t$pctAT), cdslength = sum(t$length)))  
})  
nuc_prof = do.call(rbind, nuc_prof)
```

I'm certain I'm doing something wrong to slow this down. I haven't waited for it to finish as I'm sure it can be faster. Any ideas?

data.table answer

Since I'm in an evangelizing mood ... here's what the fast `data.table` solution would look like:


```
library(data.table)
dt <- data.table(nuc, key="gene_id")

dt[,list(A=min(start),
        B=max(end),
        C=mean(pctAT),
        D=mean(pctGC),
        E=sum(length)), by=key(dt)]
```

#	gene_id	A	B	C	D	E
# 1:	NM_032291	67000042	67108547	0.5582567	0.4417433	283
# 2:	ZZZ	67000042	67108547	0.5582567	0.4417433	283

[link](#) | [edit](#) | [flag](#)

answered Jun 15 at 16:14



Josh O'Brien

20.4k ● 2 ● 14 ● 40

NB: It isn't just the speed, but the simplicity. It's easy to write and easy to read.

User's reaction

"data.table is awesome! That took about 3 seconds for the whole thing!!!"

Davy Kavanagh, 15 Jun 2012

but ...

- Example had **by=key (dt)** ?
- Yes, but it didn't need to.
- If the data is very large (1GB+) and the groups are big too then getting the groups together in memory can speed up a bit (cache efficiency).

DT[,by=] -vs- DT[,keyby=]

- **by** preserves order of groups (by order of first appearance)
- Both preserve order of rows within groups (important!) and unlike SQL
- **keyby** is a **by** as usual, followed by `setkeyv(DT, by)`

Prevailing joins (roll=TRUE)

- One reason for setkey's design.
- Last Observation (the prevailing one) Carried Forward (LOCF), efficiently
- Roll forwards or backward
- Roll the last observation forwards, or not
- Roll the first observation backwards, or not
- Limit the roll; e.g. 30 days (roll = 30)
- Join to nearest value (roll = "nearest")
- i.e. ***ordered joins***

... continued

- `roll = [-Inf, +Inf] |
 TRUE | FALSE |
 "nearest"`
- `rollends = c(FALSE, TRUE)`
- By example ...

id	date	price
SBRY	20080501	380.50
SBRY	20080502	391.50
SBRY	20080506	389.00
VOD	20080501	159.30
VOD	20080502	163.30
VOD	20080506	160.80

1. `PRC[J("SBRY")]` # all 3 rows
2. `PRC[J("SBRY",20080502),price]` # 391.50
3. `PRC[J("SBRY",20080505),price]` # NA
4. `PRC[J("SBRY",20080505),price,roll=TRUE]` # 391.50
5. `PRC[J("SBRY",20080601),price,roll=TRUE]` # 389.00
6. `PRC[J("SBRY",20080601),price,roll=TRUE,rollends=FALSE]` # NA
7. `PRC[J("SBRY",20080601),price,roll=20]` # NA
8. `PRC[J("SBRY",20080601),price,roll=40]` # 389.00

Performance

All daily prices 1986-2008 for all non-US equities

- 183,000,000 rows (id, date, price)
- 2.7 GB

`system.time(PRICES[id=="VOD"])` # vector scan

user	system	elapsed
66.431	15.395	81.831

`system.time(PRICES[J("VOD")])` # binary search

user	system	elapsed
0.003	0.000	0.002

`setkey(PRICES, id, date)` needed first (one-off apx 20 secs)

Variable name repetition

- The 3rd highest voted [R] question (of 43k)
How to sort a dataframe by column(s) in R (*)
- `DF[with(DF, order(-z, b)),]`
- VS -
`DT[order(-z, b)]`
- `quarterlyreport[with(lastquarterlyreport,order(-z,b)),]`
- VS -
`quarterlyreport[order(-z, b)]`
(*) Click link for more information

Silent incorrect results due to using a similar variable by mistake. Easily done when this appears on a page of code.

but ...

- Yes `order()` is slow when used in `i` because that's base R's `order()`.
- That's where "optimization before evaluation" comes in. We now auto convert `order()` to the internal `forder()` so you don't have to know.
- Available in v1.9.3 on GitHub, soon on CRAN

split-apply-combine

Why "split" 10GB into many small groups???

Since 2010, data.table :


- Allocates memory for largest group
- Reuses that same memory for all groups
- Allocates result data.table up front
- Implemented in C
- eval() of j within each group

Recent innovations


- Instead of the `eval(j)` from C, `dplyr` converts to an Rcpp function and calls that from C. Skipping the R eval step.
- In response, `data.table` now has **GForce**: one function call that computes the aggregate across groups. Called once only so no need to speed up many calls!
- Both approaches limited to simple aggregates: sum, mean, sd, etc. But often that's all that's needed.

data.table over-allocates

data.frame



data.table



Assigning to a subset



31



15

In R I find myself doing something like this a lot:

```
adataframe[adataframe$col==something]<-  
adataframe[adataframe$col==something])+1
```

This way is kind of long and tedious. Is there some way for me to reference the object I am trying to change such as

```
adataframe[adataframe$col==something]<-$self+1
```

?

continued

▲ Try package `data.table` and its `:=` operator. It's very fast and very short.

31

```
DT[col1==something, col2:=col3+1]
```



The first part `col1==something` is the subset. You can put anything here as if they are variables; i.e., no need to use `$`. Then the second part `col2:=col3+1` is the assignment to the LHS within that subset, where the column names can be assigned to

Easy to write, easy to read

Multiple :=

```
DT[, `:=`(newCol1=mean(colA),  
          newCol2=sd(colA)),  
    by=sector]
```

Can include a subset in i as well

set* functions

- `set()`
- `setattr()`
- `setnames()`
- `setcolorder()`
- `setkey()`
- `setkeyv()`

copy()

- data.table **IS** copied-on-change by `<-` and `=` as usual in R. Those ops aren't changed.
- No copy by `:=` or `set*`
- You have to use those, so it's clear to readers of your code
- When you need a copy, call `copy(DT)`
- Why copy a 20GB data.table, even once.
- Why copy a whole column, even once.

list columns

- Each cell can be a different type
- Each cell can be vector
- Each cell can itself be a data.table
- Combine list columns with i and by

list column example

```
data.table(  
  x = letters[1:3],  
  y = list( 1:10,  
            letters[1:4],  
            data.table(a=1:3,b=4:6)  
))
```

	x	y
1:	a	1,2,3,4,5,6,
2:	b	a,b,c,d
3:	c	<data.table>

All options

<code>datatable.verbose</code>	<code>FALSE</code>
<code>datatable.nomatch</code>	<code>NA_integer_</code>
<code>datatable.optimize</code>	<code>Inf</code>
<code>datatable.print.nrows</code>	<code>100L</code>
<code>datatable.print.topn</code>	<code>5L</code>
<code>datatable.allow.cartesian</code>	<code>FALSE</code>
<code>datatable.alloccol</code>	<code>quote(max(100L,ncol(DT)+64L))</code>
<code>datatable.integer64</code>	<code>"integer64"</code>

All symbols

- **. N**
- **. SD**
- **. I**
- **. BY**
- **. GRP**

.SD

```
stocks[, head(.SD,2), by=sector]
```

```
stocks[, lapply(.SD, sum), by=sector]
```

```
stocks[, lapply(.SD, sum), by=sector,  
.SDcols=c("mcap",paste0("revenueFQ",1:8))]
```

.I

```
if (length(err <- allocation[,  
      if(length(unique(Price))>1) .I,  
      by=stock ]$V1 )) {  
  
  warning("Fills allocated to different  
accounts at different prices! Investigate.")  
  
  print(allocation[err])  
  
} else {  
  
  cat("Ok    All fills allocated to each  
account at same price\n")  
  
}
```


Analogous to SQL

```
DT[ where,  
    select | update,  
    group by ]  
[ having ]  
[ order by ]  
[ i, j, by ] ... [ i, j, by ]
```

New in v1.9.2 on CRAN

- 37 new features and 43 bug fixes
- `set()` can now add columns just like `:=`
- `.SDcols` “de-select” columns by name or position; e.g.,

```
DT[, lapply(.SD, mean), by=colA, .SDcols=-c(3, 4)]
```
- `fread()` a subset of columns
- `fread()` commands; e.g.,

```
fread("grep blah file.txt")
```
- Speed gains

Radix sort for integer

- R's method="radix" is not actually a radix sort ... it's a counting sort. See ?setkey/Notes.
- data.table liked and used it, though.
- A true radix sort caters for range $> 100,000$
- (Negatives was a one line change to R we suggested and was accepted in R 3.1)
- Adapted to integer from Terdiman and Herf's code for float ...

Radix sort for numeric

- R reminder: numeric == floating point numbers
- Radix Sort Revisited, Pierre Terdiman, 2000
<http://codercorner.com/RadixSortRevisited.htm>
- Radix Tricks, Michael Herf, 2001
<http://stereopsis.com/radix.html>
- Their C code now in data.table with minor changes; e.g., NA/NaN and 6-pass for double

Faster for those cases

20 million rows x 4 columns, 539MB

a & b (numeric), c (integer), d (character)

	<u>v1.8.10</u>	<u>v1.8.11</u>
setkey(DT, a)	54.9s	7.2s
setkey(DT, c)	48.0s	7.0s
setkey(DT, a, b)	102.3s	16.9s

"Cold" grouping (no setkey first) :

DT[, mean(b), by=c]	47.0s	8.7s
---------------------	-------	------

<https://gist.github.com/arunsrinivasan/7997273>

New feature: melt/cast

i.e. reshape2 for data.table

20 million rows x 6 columns (a:f) 768MB

melt(**DF**, id="d", measure=1:2) 191 sec

melt(**DT**, id="d", measure=1:2) 3 sec

dcast(**DF**, d~e, ..., fun=sum) 184 sec

dcast(**DT**, d~e, ..., fun=sum) 28 sec

<https://gist.github.com/arunsrinivasan/7839891>

Similar to `melt_` in `Kmisc` by Kevin Ushey

... melt/cast continued

Q: Why not submit a pull request to reshape2 ?

A: This C implementation calls data.table internals at C-level (e.g. fastorder, grouping, and joins). It makes sense for this code to be together.

Miscellaneous 1

```
DT[, (myvar) := NULL]
```

Spaces and specials; e.g., `by="a, b, c"`

```
DT[4:7, newCol := 8] []
```

- extra `[]` to print at prompt
- auto fills rows 1:3 with NA

```
rbindlist(lapply(fileNames, fread))
```

`fread`'s `drop` and `select`

Miscellaneous 2

Dates and times

Errors & warnings are deliberately very long

Not joins `X [! Y]`

Column plonk & non-coercion on assign

by-without-by \Rightarrow `by=.EACHI`

Secondary keys / merge

R3, singleton logicals, reference counting

`bit64::integer64`

Miscellaneous 3

Print method vs typing DF, copy fixed in R-devel

How to benchmark

mult = "all" | "first" | "last" (may expand)

with=FALSE

which=TRUE

nomatch 0 (inner join) or NA (outer join)

Chained queries: DT[...][...][...]

Dynamic and flexible queries (eval text and quote)

fread sep2

Not (that) much to learn

- Main manual page: `?data.table`
- Run `example(data.table)` at the prompt (53 examples)
- No methods, no functions, just use what you're used to in R

Thank you

<https://github.com/Rdatatable/datatable/>

<http://stackoverflow.com/questions/tagged/data.table>

```
> install.packages("data.table")  
> require(data.table)  
> ?data.table  
> ?fread
```

Learn by example :

```
> example(data.table)
```